

Self-Learning Stock Trading Bot

- 1.) Ananya Shroff- UCE2023563
- 2.) Divyanshi Singh - UCE2023565
- 3.) Sanika Tavate - UCE2023566
- 4.) Trupti Patil-UEC2023I63





Introduction

- **The Challenge:** Financial markets are highly volatile and complex. Developing automated trading strategies that can adapt to changing market conditions and manage risk is extremely difficult.
-
- **The Goal:** To design and implement an intelligent agent that learns optimal, profitable trading strategies directly from market data through trial and error.
-
- **Our Project:** We present an end-to-end system featuring a Deep Q-Network (DQN), a type of reinforcement learning agent, trained for automated stock trading.
-
- **Key Feature:** The system includes a dynamic, real-time web-based visualization interface to simulate, analyze, and monitor the agent's performance interactively.





Our Solution




We developed an integrated system with three main functions:

1. **Automated Learning:** A reinforcement learning agent (DQN) that trains on historical stock data to discover complex trading patterns and strategies on its own.
2. **Integrated Risk Management:** The trading environment is built with practical, automatic risk controls, including Stop-Loss (5%) and Take-Profit (8%) triggers, position size limits, and transaction cost modeling (0.1% per trade).
3. **Real-Time Visualization:** A web-based dashboard that simulates the agent's performance, streaming every trade, decision, and portfolio change live to the user for in-depth analysis.



AIML Techniques Used

01 **Reinforcement Learning (RL)**: The core methodology where an "agent" learns to make optimal decisions (actions) in an "environment" to maximize a cumulative "reward."



02 **Deep Learning**: Using a multi-layer neural network to approximate the complex Q-value function, which estimates the expected future reward for each possible action.

03 • **Deep Q-Network (DQN)**: The specific algorithm that combines Reinforcement Learning with a Deep Neural Network. It uses key features like:


- Experience Replay Buffer: Stores past experiences (state, action, reward, next state) and samples from them randomly to break correlations and stabilize learning.
- Target Network: A separate, periodically updated copy of the main network used to create more stable learning targets.

04 • **Advanced Feature Engineering**:

- Temporal Lookback: A 30-timestep rolling window of historical returns, allowing the agent to see recent price momentum.
- Technical Indicators: 13 distinct, normalized indicators (RSI, MACD, Bollinger Bands, etc.) provide the agent with a rich understanding of the market state.
- Data Normalization: Using Z-scores and tanh scaling to ensure all features are in a stable range for the neural network.



DQN Model Architecture



Input Layer

43 features for market analysis

Hidden Layers

2 layers with 256 neurons each

Output Layer

3 Q-values for action selection

Action Space


3 discrete actions: hold, buy, sell

Observation Space

43-D vector of market data
30 past normalized returns,
13 normalized features

Reward Function

Maximizes portfolio value



Tech Stack

1. Backend & AI:

- Python 3.10+: Core programming language.
- Stable-Baselines3: The library used for the DQN algorithm.
- PyTorch: The deep learning framework powering the neural network.
- Gymnasium (OpenAI Gym): The toolkit used to build the custom trading environment.
- FastAPI: High-performance web framework for the backend API.
- WebSockets: For real-time, bidirectional communication with the frontend.
- Pandas & NumPy: For data manipulation and numerical operations.


2. Frontend (Visualization):

- HTML5, CSS3, JavaScript: The foundation of the web dashboard.
- Plotly.js: The interactive charting library used to render all graphs and charts.

3. Data:

- yfinance: Python library to download historical stock data from Yahoo Finance.
- CSV: Storage format for processed data.



Project Components



Made with DataCamp



Real-Time Dashboard Components



Trading Bot Simulator

AAPL ▾ Speed 2.0s Use Model

Connected

Play

Pause

Stop

Price Chart



Portfolio

Net Worth

\$109024.23

Cash

\$21580.53

Unrealized P/L

\$9024.23

Value (\$)

\$100,000

\$105,000

\$110,000

\$105,000

\$100,000

\$95,000

\$90,000

\$85,000

\$80,000

\$75,000

\$70,000

\$65,000

\$60,000

\$55,000

\$50,000

\$45,000

\$40,000

\$35,000

\$30,000

\$25,000

\$20,000

\$15,000

\$10,000

\$5,000

\$0,000

Aug 29 2021 Sep 12 Sep 26 Oct 10 Oct 24 Nov 7 Nov 21

Time

Recent Actions

BUY 265 shares @ \$161.94

BUY 264 shares @ \$161.41

SELL 133 shares @ \$161.02

BUY 267 shares @ \$148.64

RIIV 267 shares @ €148.60

Statistics

Total Return

9.02%

Current Step

98

Cumulative Reward

7.49

System Log

[11:34:25 PM] Disconnected

[11:34:19 PM] BUY 265 shares @ \$161.94

[11:34:16 PM] BUY 264 shares @ \$161.41

[11:34:14 PM] SELL 133 shares @ \$161.02

[11:33:34 PM] BUY 267 shares @ \$148.64

[11:33:22 PM] RIIV 267 shares @ €148.60

Disconnected

*

*

*

*

*

*

*

*

*

*

*

*

*

*

*

*

*

*

*

*

*

*

*

*

*

*

*

*

*

*

*

*

*

*

*

*

*

*

*

*

*

*

*

*

*

*

*

*

*

*

*

*

*

*

*

*

*

*

*

*

*

*

*

*

*

*

*

*

*

*

*

*

*

*

*

*

*

*

*

*

*

*

*

*

*

*

*

*

*

*

*

*

*

*

*

*

*

*

*

*

*

*

*

*

*

*

*

*

*

*

*

*

*

*

*

*

*

*

*

*

*

*

*

*

*

*

*

*

*

*

*

*

*

*

*

*

*

*

*

*

*

*

*

*

*

*

*

*

*

*

*

*

*

*

*

*

*

*


*

*

*

System Architecture

Data Pipeline Process



01
Download Raw Data

Fetches OHLCV data from Yahoo Finance

02
Calculate Indicators

Computes 13 technical indicators and returns

03
Normalize Data

Scales data to a standard range

04
Save to CSV


Stores processed data in a CSV file

Trading Environment Setup

Load Processed CSV

Define Reset Function


Define Observation Spaces



Define Step Function

Define Action Spaces
Define Reward Function

DQN Model Training Workflow



Made with ➜ Napkin

Web Simulator Process



Made with ➜ Napkin

Conclusion



We successfully built and demonstrated a complete, end-to-end system for developing and evaluating reinforcement learning-based trading agents.

- The Deep Q-Network model, combined with a rich feature set and integrated risk management, forms a robust foundation for automated trading.
- The real-time, interactive web dashboard proves to be an invaluable tool for analyzing agent behavior and building trust in the system.

- **Current Limitations:**

- Single Asset: The agent can only trade one stock at a time.
- Fixed Position Size: The agent can only Buy/Sell 1 share, not decide how much to trade.
- Offline Only: The system runs on historical data, not live market data.

- **Future Work:**

- Portfolio Management: Expand the agent to trade multiple assets simultaneously.
- Variable Position Sizing: Allow the agent to learn the optimal amount of capital to risk per trade.
- Live Trading Integration: Connect the system to a brokerage API for real-time paper and live trading.
- Explore Other Models: Compare DQN's performance against other modern RL algorithms like PPO or A3C.
- Alternative Features: Integrate new data sources, such as market news sentiment.

..... References

- Mnih, V., et al. (2015). "Human-level control through deep reinforcement learning." *Nature*.
- Jiang, Z., et al. (2017). "A Deep Reinforcement Learning Framework for the Financial Portfolio Management Problem." *arXiv*.
- Yang, H., et al. (2020). "Deep Reinforcement Learning for Stock Trading." *IEEE*.
- Raffin, A., et al. (2021). "Stable-Baselines3: Reliable Reinforcement Learning Implementations." *Journal of Machine Learning Research (JMLR)*.



.....



Thank You